

Linear Regression

[Code ▾](#)

Import Library

[Hide](#)

```
library(Metrics)
library(DAAG)
```

Get the data set

[Hide](#)

```
data_set <- read.csv("Advertising.csv", header=TRUE)
```

Summary of the data set

[Hide](#)

```
summary(data_set)
```

X	TV	Radio	Newspaper	Sales
Min. : 1.00	Min. : 0.70	Min. : 0.000	Min. : 0.30	Min. : 1.60
1st Qu.: 50.75	1st Qu.: 74.38	1st Qu.: 9.975	1st Qu.: 12.75	1st Qu.: 10.38
Median : 100.50	Median : 149.75	Median : 22.900	Median : 25.75	Median : 12.90
Mean : 100.50	Mean : 147.04	Mean : 23.264	Mean : 30.55	Mean : 14.02
3rd Qu.: 150.25	3rd Qu.: 218.82	3rd Qu.: 36.525	3rd Qu.: 45.10	3rd Qu.: 17.40
Max. : 200.00	Max. : 296.40	Max. : 49.600	Max. : 114.00	Max. : 27.00

[Hide](#)

```
print(nrow(data_set))
```

```
[1] 200
```

[Hide](#)

```
print(ncol(data_set))
```

```
[1] 5
```

Drop index column

[Hide](#)

```
data_set <- data_set[c(2:5)]  
  
colnames(data_set)
```

```
[1] "TV"      "Radio"    "Newspaper" "Sales"
```

SPLIT DATA SET INTO TRAINING AND TESTING

80% of the sample size

[Hide](#)

```
ind_split <- floor(0.80 * nrow(data_set))
```

Set seed to make your partition reproducible

[Hide](#)

```
set.seed(123)  
train_ind <- sample(seq_len(nrow(data_set)), size = ind_split)  
  
train <- data_set[train_ind, ]  
test <- data_set[-train_ind, ]  
  
View(train)  
View(test)
```

LINEAR REGRESSION

[Hide](#)

```
lr_age <- lm(Sales~TV, data=train)  
summary(lr_age)
```

```
Call:
lm(formula = Sales ~ TV, data = train)

Residuals:
    Min       1Q   Median       3Q      Max
-8.3396 -1.9922  0.0219  2.0201  7.2355

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  7.069052   0.507788   13.92  <2e-16 ***
TV           0.047237   0.002983   15.83  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.257 on 158 degrees of freedom
Multiple R-squared:  0.6134,    Adjusted R-squared:  0.611
F-statistic: 250.7 on 1 and 158 DF,  p-value: < 2.2e-16
```

Hide

```
lr_dist <- lm(Sales~Radio, data=train)
summary(lr_dist)
```

```
Call:
lm(formula = Sales ~ Radio, data = train)

Residuals:
    Min       1Q   Median       3Q      Max
-15.5706 -2.1136  0.8746  2.9348  8.3634

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  9.41474    0.64372   14.626  < 2e-16 ***
Radio        0.19586    0.02325    8.426 2.08e-14 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.352 on 158 degrees of freedom
Multiple R-squared:  0.31, Adjusted R-squared:  0.3056
F-statistic: 70.99 on 1 and 158 DF,  p-value: 2.083e-14
```

Hide

```
lr_conv <- lm(Sales~Newspaper, data=train)
summary(lr_conv)
```

```
Call:
lm(formula = Sales ~ Newspaper, data = train)

Residuals:
    Min       1Q   Median       3Q      Max
-11.728  -3.717  -1.147   3.767  12.643


Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) 13.05783    0.70700  18.469  <2e-16 ***
Newspaper    0.03109    0.01901   1.635   0.104
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.195 on 158 degrees of freedom
Multiple R-squared:  0.01664,    Adjusted R-squared:  0.01042
F-statistic: 2.674 on 1 and 158 DF,  p-value: 0.104
```


PLOT GRAPH FOR RELATIONSHIP BETWEEN FEATURES AND TARGET

[Hide](#)


```
plot(train$Sales~train$TV, xlab="TV", ylab ="Sales")
abline(lr_age)
```

[Hide](#)

```
plot(train$Sales~train$Radio, xlab="Radio", ylab = "Sales")  
abline(lr_dist)
```

[Hide](#)

```
plot(train$Sales~train$Newspaper, xlab="Newspaper", ylab ="Sales")  
abline(lr_conv)
```



Calculate MSE for distance

[Hide](#)

```
lr_dist_1 <- lm(Sales~log(TV), data=train)
View(lr_dist_1)
```

Prediction on train

[Hide](#)

```
lr_train_1 <- predict(lr_dist_1, train)
lr_train_1
```

	159	179	14	195	170	50	118	43
198	194	153						
	5.643334	17.687556	13.716094	15.348648	17.790723	12.282012	12.787577	17.913278
6456	15.760469	16.405635						15.98
	90	91	188	185	92	137	99	72
26	7	196						
	14.168448	14.935324	16.278285	17.358643	9.046485	8.624567	17.862363	14.168448
2768	11.705513	10.148466						17.49
	184	164	78	81	193	103	117	76
143	32	109						
	17.834663	15.684387	14.522498	12.787577	7.110416	17.735415	15.071766	7.043421
3131	14.274454	6.073663						16.82
	180	178	74	23	155	53	135	162
163	34	69						
	15.732978	15.837298	14.793811	6.102617	16.211962	16.751669	10.016637	13.224937
4107	17.531671	17.104306						16.22
	182	171	63	141	97	187	38	21
41	189	60						
	16.788439	11.173379	17.134657	12.635056	16.405635	15.079963	12.701900	16.786696
8899	17.813422	16.650037						16.49
	16	116	94	6	86	190	39	191
197	158	4						
	16.363007	12.722233	17.314887	4.515322	16.319896	7.428771	10.607975	10.275882
4996	15.351191	15.394156						13.58
	13	157	127	52	22	89	169	110
25	87	35						
	8.346979	13.572851	4.099553	13.827689	17.104306	13.338731	16.734034	17.382570
0779	12.782590	13.645146						12.01
	40	112	30	12	31	132	121	64
183	192	93						
	16.950482	17.172652	12.486970	16.721641	17.904189	17.525932	15.128777	13.913927
8444	12.742458	16.774474						11.61
	96	71	67	177	79	85	37	8
51	165	166						
	15.679726	16.434429	9.414210	17.276759	2.699465	16.700300	17.550261	14.513007
7792	14.416774	17.057509						16.44
	98	173	140	200	84	46	17	62
122	54	124						
	16.152709	7.607743	16.152709	17.018341	12.366437	15.945364	12.332891	17.469525
9077	16.105051	14.603777						7.44
	108	24	125	167	147	168	181	128
88	27	42						
	13.428221	16.955489	16.975449	7.262299	17.147364	16.578902	15.520217	12.972393
9529	15.171648	15.986456						14.19
	5	70	145	129	83	149	55	36
139	186	9						
	16.067332	16.758700	13.664987	16.819676	12.732359	10.128479	17.489870	17.875483
9131	16.545616	4.471305						10.59
	48	56	111	1	123	172	154	131
126	77	101						


```

17.144191 16.430602 16.913565 16.985390 16.883092 15.707603 15.861826 -5.079401 13.29
1002 8.897155 16.855799
      176      11      68      20      144      3      29      156
44      130      80
17.690307 12.236207 15.074500 15.287113 13.983723 7.110416 17.282885 1.650853 16.58
0742 11.842088 14.377589
      57      138      105      100      2      134
3.847313 17.646051 17.117115 14.960754 10.729684 16.811025

```

Prediction on test

[Hide](#)

```

lr_pred_1 <- predict(lr_dist_1, test)
lr_pred_1

```

```

      10      15      18      19      28      33      45      47
49      58      59
16.447792 16.528864 17.751686 12.410710 17.147364 13.704361 8.549467 13.398624 16.93
7099 14.988812 16.651843
      61      65      66      73      75      82      95      102
104      106      107
11.430985 14.843505 12.399690 8.798983 16.698517 17.142604 14.084302 17.949416 16.21
3989 15.036041 8.534268
      113      114      115      119      120      133      136      142
146      148      150
15.958389 16.630107 12.876240 14.683356 7.568692 4.381714 11.041674 16.329737 15.10
1735 17.196208 10.746758
      151      152      160      161      174      175      199
17.742203 14.538264 14.860891 15.888405 15.796817 16.855799 17.781337

```

Calculate MSE

[Hide](#)

```

train_mse_1 = mse(train$Sales, lr_train_1)
print(train_mse_1)

```

```
[1] 11.48834
```

[Hide](#)

```

test_mse_1 = mse(test$Sales, lr_pred_1)
print(test_mse_1)

```

```
[1] 13.00607
```

Graph of test vs train mse

[Hide](#)

```
plotter <- c(train_mse_1, test_mse_1)
barplot(plotter, width = 0.02, xlab="data", names.arg = c("Train MSE", "Test MSE"), ylab="error", main="Error (TV Sales)")
```



Correlation (Subset Selection Method)

[Hide](#)

```
res <- cor(train)
print(res)
```

```

      TV      Radio  Newspaper   Sales
TV      1.00000000 0.04039595 -0.03283919 0.783218
Radio    0.04039595 1.00000000  0.33173584 0.556789
Newspaper -0.03283919 0.33173584  1.00000000 0.128995
Sales     0.78321799 0.55678901  0.12899496 1.000000
```

[Hide](#)

```
lin_gen <- lm(Sales~TV+Radio+Newspaper, data=train)
print(lin_gen)
```

Call:

```
lm(formula = Sales ~ TV + Radio + Newspaper, data = train)
```

Coefficients:

(Intercept)	TV	Radio	Newspaper
3.04081	0.04589	0.18774	-0.00556

Prediction on train

[Hide](#)

```
lin_train <- predict(lin_gen, train)
lin_train
```

159	179	14	195	170	50	118	43
198	194	153					
10.253887	16.038942	8.901963	16.560754	18.042100	8.102846	6.614804	21.704741
3914	18.560345	16.404208					
90	91	188	185	92	137	99	72
26	7	196					
16.767666	10.072199	17.097449	18.519984	4.451435	11.485607	23.992065	10.588046
4275	11.706631	5.411751					
184	164	78	81	193	103	117	76
143	32	109					
23.912583	17.411552	13.842219	11.435460	4.424171	17.676685	11.971161	11.523394
1939	11.273929	3.574752					
180	178	74	23	155	53	135	162
163	34	69					
12.419903	12.120137	9.875212	6.355815	15.567585	20.580087	11.616083	13.420535
2405	18.982561	19.036959					
182	171	63	141	97	187	38	21
41	189	60					
13.929482	7.410806	16.780705	9.529010	12.733187	9.688995	15.488950	17.966850
4604	18.754653	18.196585					
16	116	94	6	86	190	39	191
197	158	4					
20.668843	12.764998	21.005313	12.203368	14.996068	6.040480	9.836120	12.537198
8632	10.024273	17.421577					
13	157	127	52	22	89	169	110
25	87	35					
10.356165	15.235736	10.420361	9.430533	14.762188	11.472192	17.036110	19.780937
3559	11.616084	7.654286					
40	112	30	12	31	132	121	64
183	192	93					
20.403735	21.137689	9.057669	17.377070	21.555095	15.516535	14.299702	13.264109
4869	8.499781	18.992459					
96	71	67	177	79	85	37	8
51	165	166					
16.173191	17.707307	9.092445	19.996945	8.849654	20.723304	23.484198	12.172060
9496	11.148950	13.969155					
98	173	140	200	84	46	17	62
122	54	124					
15.346233	7.619248	19.758237	15.258319	14.336082	15.125287	12.389571	22.744390
7225	19.767587	15.116731					
108	24	125	167	147	168	181	128
88	27	42					
7.116713	16.544847	19.224188	10.801032	15.381393	13.399484	10.669342	6.670135
1670	15.029266	17.218770					
5	70	145	129	83	149	55	36
139	186	9					
13.040810	21.080400	10.017765	22.331899	10.126764	12.284259	20.414822	17.103847
2511	20.806418	3.824156					
48	56	111	1	123	172	154	131
126	77	101					

```

21.738289 21.109140 14.628366 20.312067 13.684279 14.250067 18.145484 10.458892 9.11
3810 4.488110 13.777416
      176      11      68      20      144      3      29      156
44      130      80
24.695972 7.028543 12.098914 14.181279 8.719875 12.061925 19.418867 5.375004 13.96
5910 7.789136 9.681328
      57      138      105      100      2      134
8.421019 20.694893 20.381963 16.818694 12.210241 19.166108

```

Prediction on test

[Hide](#)

```

lin_pred <- predict(lin_gen, test)
lin_pred

```

```

      10      15      18      19      28      33      45      47
49      58      59
12.580125 18.327971 23.078708 9.963320 17.067163 7.616252 8.776754 8.817358 16.15
6097 12.803428 21.816770
      61      65      66      73      75      82      95      102
104      106      107
5.752485 16.931561 7.948238 10.358667 17.379461 14.610094 10.537224 22.896816 14.79
3306 17.752128 6.088061
      113      114      115      119      120      133      136      142
146      148      150
13.981681 16.467475 15.223724 15.296465 6.810888 8.521026 14.033670 18.155479 9.78
6008 23.154312 9.821197
      151      152      160      161      174      175      199
18.326298 9.899876 12.346657 14.184381 12.030651 13.812488 23.572416

```

Calculate MSE on subset

[Hide](#)

```

train_mse = mse(train$Sales, lin_train)
print(train_mse)

```

```
[1] 2.978008
```

[Hide](#)

```

test_mse = mse(test$Sales, lin_pred)
print(test_mse)

```

```
[1] 2.068075
```

Graph of test vs train mse for subset

[Hide](#)

```
plotter <- c(train_mse, test_mse)
barplot(plotter, width = 0.02, xlab="data", names.arg = c("Train MSE", "Test MSE"), ylab="error", main="Error (Subset)")
```



K fold cross validation

[Hide](#)

```
model = cv.lm(df, (Sales~TV+Radio+Newspaper), m=5)
```